# Kathmandu University

# Department of Computer Science and Engineering

# Dhulikhel, Kavre



**A Project Report**

**on**

**"World's Hardest Game Solving Bot"**

**[Code No: COMP 308]**

**(For partial fulfillment of Third Year/ Second Semester in Computer Engineering)**

**Submitted by:**

**Aakriti Banjara (05)**

**Jayash Bhattarai (08)**

**Regal Adhikari (64)**

**Submitted to**

**Mr. Amrit Dahal**

**Department of Computer Science and Engineering**

**Submission Date: 21th June 2024**

# Bona fide Certificate

**This project work on**

**"World's Hardest Game Solving Bot"**

**is the bona fide work of**

**"<u>Aakriti Banjara, Jayash Bhattarai and Regal Adhikari</u>"**

**who carried out the project work under my supervision.**

**Project Supervisor**

_____

**Dr Rajani Chulyadyo**

**Assistant Professor,**

**Department of Computer Science and Engineering (DoCSE)**

**Date: 21th June 2024**

# ACKNOWLEDGEMENT

# Abstract

The "World's Hardest Game" project innovatively combines game design with advanced AI techniques, specifically Q-learning reinforcement learning, to deliver an engaging experience across three challenging levels. Players navigate a character through intricate mazes, evading enemies and obstacles to reach designated endpoints. The project features a user-playable mode with keyboard controls and an autonomous bot mode where a neural network-based reinforcement learning model learns optimal strategies through trial and error, maximizing rewards and minimizing penalties. Developed using Python and Pygame, the game offers an immersive graphical interface and showcases the practical application of reinforcement learning in gaming. This project demonstrates how AI can master complex tasks in dynamic environments, contributing to advancements in AI-driven gaming systems and enhancing player experience through adaptive opponent behavior. The report outlines methodologies, achievements, challenges, and future enhancement recommendations, highlighting its contributions to gaming technology and AI research.

**Keywords**: Reinforcement learning, Q-learning Algorithm, Pygame

# Table of Contents

# List of Figures

# Acronyms/Abbreviations

- RL: Reinforcement Learning

- Q-learning: Q-learning is a model-free reinforcement learning algorithm. The algorithm's goal is to find the best action to take based on the current state, and it aims to maximize the expected reward over all the successive steps.

- DQN: The integration of artificial Neural Nets (NNs) into the Q-learning process is referred to as Deep Q-learning, and a network that uses NNs to approximate Q-functions is called a Deep Q-Network (or DQN).

# Chapter 1: Introduction

The "World's Hardest Game" project aims to provide players with a challenging gaming experience across three uniquely designed levels. The game challenges players to navigate a character from a starting point to a destination while avoiding enemies and obstacles strategically placed throughout each level. In addition to manual gameplay, the project implements a reinforcement learning model that learns to play the game autonomously through trial and error, optimizing its actions based on rewards and penalties received during gameplay.

## 1.1 Background

The advent of artificial intelligence (AI) and machine learning (ML) has revolutionized various industries, including gaming. AI techniques, particularly reinforcement learning (RL), have garnered significant attention for their ability to enable autonomous decision-making in dynamic and complex environments. Reinforcement learning (RL) is a machine learning (ML) technique that trains software to make decisions to achieve the most optimal results. It mimics the trial-and-error learning process that humans use to achieve their goals. The application of RL in gaming not only enhances player experience but also pushes the boundaries of AI capabilities in strategic planning, problem-solving, and adaptive learning.

## 1.2 Problem Statement

The "World's Hardest Game" serves as a challenging benchmark in gaming, requiring players to navigate through intricate mazes while avoiding enemies to reach designated goals. While players can manually control their avatars, developing an AI agent capable of mastering such gameplay requires overcoming numerous challenges. These include formulating efficient learning algorithms, designing robust state representations, and implementing effective reward mechanisms to incentivize desired behaviors.

## 1.3 Motivation

This project aims to integrate game development with deep learning reinforcement learning. The motivation behind this project comes from the desire to create a gaming experience with an AI bot learning its way to complete through the levels and finally reaching to the end. As the game gets

tougher and tougher with increasing levels, the AI analyzes the situation and slowly gets its way to complete each of them. This project also contributes in the field of AI-driven gaming which has abundant potential in the gaming industry as well as the AI industry of recent times.

## 1.4 Objectives

The primary objective of this project is to develop a deep reinforcement learning agent capable of autonomously playing the "World's Hardest Game" Specific objectives include:

1. Implementing RL Algorithms: Deploying Q-learning and Deep Q-Networks (DQNs) to facilitate autonomous decision-making.
2. Navigating Complex Environments: Designing strategies for the agent to navigate mazes, avoid enemies, and reach goals efficiently.
3. Achieving Competitive Performance: Benchmarking the agent's performance against human players and established gaming benchmarks.
4. Contributing to AI Research: Contributing insights into the application of RL in gaming contexts and advancing the field of autonomous gameplay.

## 1.5 Scope of Work

This project focuses on developing an RL-based agent for three distinct levels of the "World's Hardest Game." The agent's performance metrics include score accumulation, completion rates, and decision-making efficiency. The scope encompasses:

1. Utilizing Q-learning and DQNs within a Python-based gaming environment.
2. Assessing the agent's performance through empirical analysis, comparing against human and baseline benchmarks.

## 1.6 Significance of the Study

This study contributes to the burgeoning field of AI and gaming by demonstrating the feasibility and efficacy of RL techniques in tackling complex gameplay challenges. Insights gained from this project can inform the development of AI-driven solutions not only in gaming but also in real-world applications requiring adaptive decision-making and strategic planning.

# Chapter 2 Related Works

A brief about technical background and related works in this field have been debriefed below:

## 2.1 Technical Background

This section provides an overview of the technical foundations that underpin the development of the AI agent, focusing on reinforcement learning, Q-learning, neural networks, and the overall system architecture.

### 2.1.1 Reinforcement Learning Framework

Reinforcement Learning (RL) is a type of machine learning where an agent learns to make decisions by performing actions in an environment to maximize cumulative reward. The RL framework consists of several key components and processes, which work together to enable the agent to learn optimal behaviors through trial and error. (bajaj, 2023)

Key Components are as follows:

- Agent: The decision-maker that interacts with the environment.
- Environment: The external system the agent interacts with, which responds to the agent's actions.
- State (s): A representation of the current situation of the agent within the environment.
- Action (a): The set of all possible moves the agent can make.
- Reward (r): A scalar feedback signal received after performing an action, indicating the immediate benefit of that action.
- Policy (π): A strategy used by the agent to determine the next action based on the current state.
- Value Function (V): A function that estimates the expected cumulative reward from a given state, under a particular policy.
- Q-Function (Q): A function that estimates the expected cumulative reward of taking a specific action in a specific state and following a particular policy thereafter.

### 2.1.2 Q-Learning Algorithm

The core learning algorithm employed is Q-learning, a model-free RL technique that learns optimal policies by iteratively updating Q-values for state-action pairs. The Q-values represent the expected cumulative reward for taking a particular action in each state, guiding the agent towards actions that maximize long-term rewards. (Chanda, 2024)

$$Q^{new}(S_t, A_t) \leftarrow (1 - \underbrace{\alpha}_{\text{learning rate}}) \cdot \underbrace{Q(S_t, A_t)}_{\text{current value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left( \underbrace{R_{t+1}}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(S_{t+1}, a)}_{\text{estimate of optimal future value}} \right)}_{\text{new value (temporal difference target)}}$$

*Figure 2.1: Q- learning algorithm*

where,

- $\alpha$ is the learning rate
- $\gamma$ is the discount factor
- $r$ is the reward received after taking action
- s' is the next state
- $a'$ is the action in the next state

The agent's policy is derived from the Q-values, typically choosing the action with the highest Q-value for the current state.

### 2.1.3 Experience Replay

Experience Replay is employed to enhance training stability and efficiency. It involves storing agent experiences (state, action, reward, next state) in a replay memory buffer. During training, random batches of experiences are sampled from this buffer, breaking the sequential correlation of experiences and improving learning efficiency.

### 2.1.4 Exploration-Exploitation Strategy

To balance exploration of new strategies with exploitation of learned behaviors, an ε-greedy strategy is adopted. At each time step, the agent either selects a random action with probability ε (exploration) or chooses the action with the highest Q-value from the current state (exploitation). The exploration rate ε decays over time, promoting more exploitation as training progresses.

## 2.2 Reinforcement Learning in Gaming

Reinforcement learning (RL) has garnered significant attention in the gaming industry due to its ability to create intelligent agents capable of learning and adapting to complex environments. Various studies have explored the application of RL algorithms in different game genres, ranging from classic board games like Chess and Go to modern video games requiring real-time decision-making and strategy like Dota 2 and VALORANT.

### 2.2.1 Deep Q-Networks (DQN)

Deep Q-Networks (DQN) pioneered the use of deep neural networks in RL, demonstrating superior performance in learning policies for navigating game environments. The integration of convolutional neural networks (CNNs) in DQN architectures, as seen in the seminal work by Mnih et al. (2015), enabled agents to learn directly from raw pixel inputs, achieving human-level performance in Atari 2600 games.

### 2.2.2 AlphaGo and AlphaStar

Google DeepMind's AlphaGo (AlphaGo, n.d.) and AlphaStar (team, 2019) projects represent milestones in RL applied to strategic board games (Go) and real-time strategy games (StarCraft II). These projects showcased RL's capability to master games with vast state and action spaces, leveraging advanced techniques such as Monte Carlo Tree Search (MCTS) and policy gradient methods.

### 2.2.3 AI in Platformer Games

Platformer games, characterized by jumping between platforms and overcoming obstacles, present unique challenges for AI agents. Prior research has explored methods for teaching agents to navigate platformer environments, focusing on pathfinding algorithms, heuristic-based decision-making, and physics-based simulations.

### 2.2.4 Mario AI Competition

The Mario AI competition, initiated by the University of Texas at Austin, encouraged research into AI agents capable of autonomously playing Super Mario Bros. Agents were evaluated based

on their ability to complete levels while maximizing score and minimizing risks, stimulating innovations in navigation and enemy interaction strategies. (Julian Togelius, 2013)

### 2.2.5 Procedural Content Generation

Research on procedural content generation (PCG) has influenced level design in games, enabling dynamic and unpredictable gameplay experiences. PCG techniques include cellular automata, genetic algorithms, and reinforcement learning-based approaches for generating challenging yet fair game levels. (Oscar, 2023)

# Chapter 3: Methodology

The methodology for applying reinforcement learning (RL) to the "World's Hardest Game" involves several key steps, including problem formulation, environment setup, model architecture, training process, and evaluation metrics. This section outlines the detailed methodology used to develop a bot capable of playing the game using deep reinforcement learning, specifically leveraging the Q-learning algorithm. It is a challenging puzzle and strategy game where the player's objective is to navigate a small square (the player) through a series of levels filled with moving enemies and obstacles. The player must reach the designated goal area without colliding with any enemies.

## 3.1 Problem Statement

The project aims to develop a reinforcement learning (RL) agent capable of autonomously playing the "World's Hardest Game." This involves creating an intelligent agent that can navigate through three distinct levels, avoiding enemies and reaching designated goals. The agent must learn optimal strategies through interaction with the game environment, leveraging the Q-learning algorithm integrated with deep neural networks.

## 3.2 Overview of the "World's Hardest Game"

The "World's Hardest Game" is a challenging puzzle and strategy game where the player's objective is to navigate a small square (the player) through a series of levels filled with moving enemies and obstacles. The player must reach the designated goal area without colliding with any enemies. The game is known for its difficulty, requiring precise movements and strategic planning to succeed. The game consists of multiple levels, each increasing in difficulty. Each level has a unique layout, including different patterns of enemy movements and obstacles. The player controls a small square that can move in four directions: Up, Down, Left, and Right. Movement must be carefully timed and planned to avoid enemies and reach the goal. Enemies are typically represented as moving circles that follow predefined paths. Colliding with an enemy results in an immediate level reset. Each level has a start point and a goal area that the player must reach to complete the level. The primary objective is to navigate from the start point to the goal area without being hit by enemies.

*Figure 3.1: level 1 Diagram*



*Figure 3.2: Level 2 Diagram*

### 3.2.1 Game Environment

Our game environment comprises of three levels, each presenting unique layouts of obstacles, enemies, and goals. The agent interacts with the game through actions such as movement and decision-making, aiming to maximize its score by safely navigating to the goal while avoiding collisions with enemies.

This flow chart outlines the sequential steps and decision points involved in the gameplay and AI training process for the "World's Hardest Game" project. The chart illustrates the following key stages:
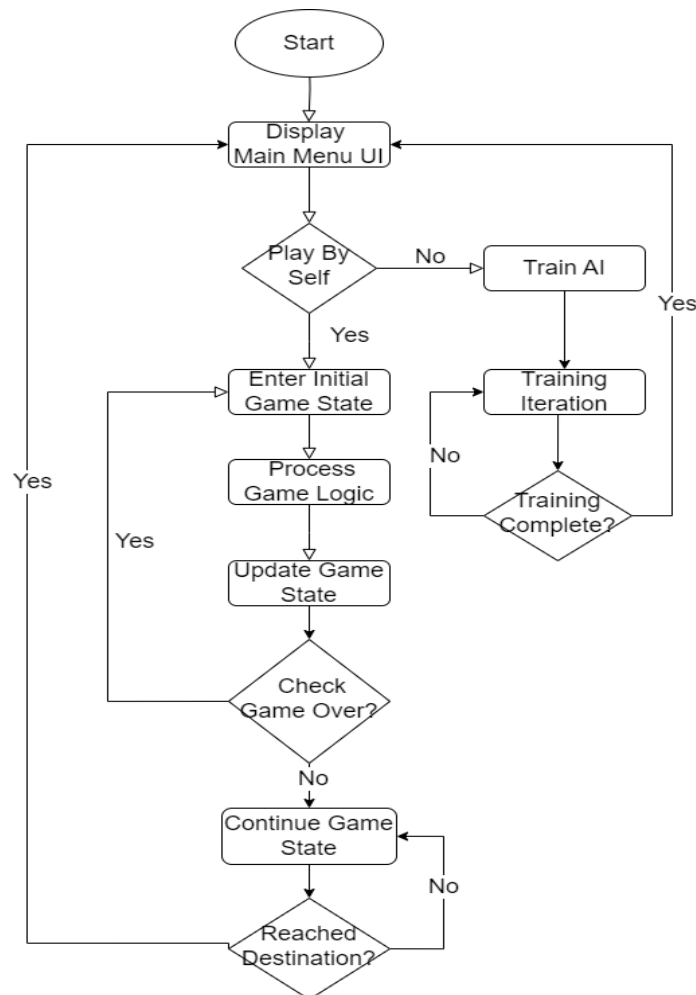


*Figure 3.3: Flow Chart*

### 3.2.2 State Representation

The state of the game environment is represented using a feature vector derived from the game's current state. This vector includes information such as the agent's position, enemy positions, and proximity to goals and obstacles. Each state is discretized to facilitate learning by the RL agent.

This state flow diagram illustrates the various states and transitions of the AI Agent as it navigates through the "World's Hardest Game". It provides a visual representation of the following elements:
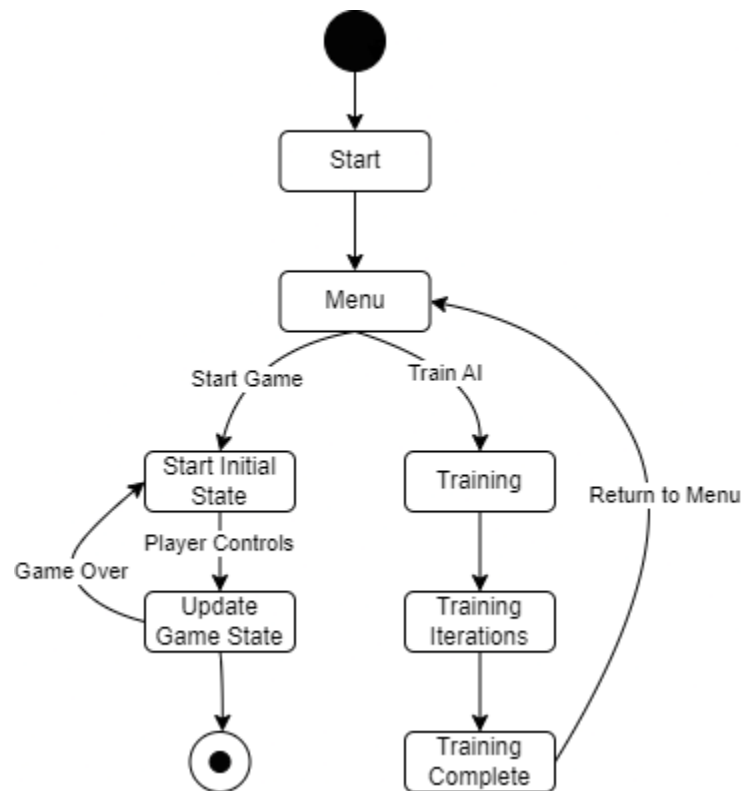


*Figure 3.4: State Flow Diagram*

### 3.2.3 Use Case Diagram

This use case diagram illustrates the interactions between the primary actors (Player and AI Agent) and the system components in the "World's Hardest Game" project. The diagram highlights the main functionalities of the system.
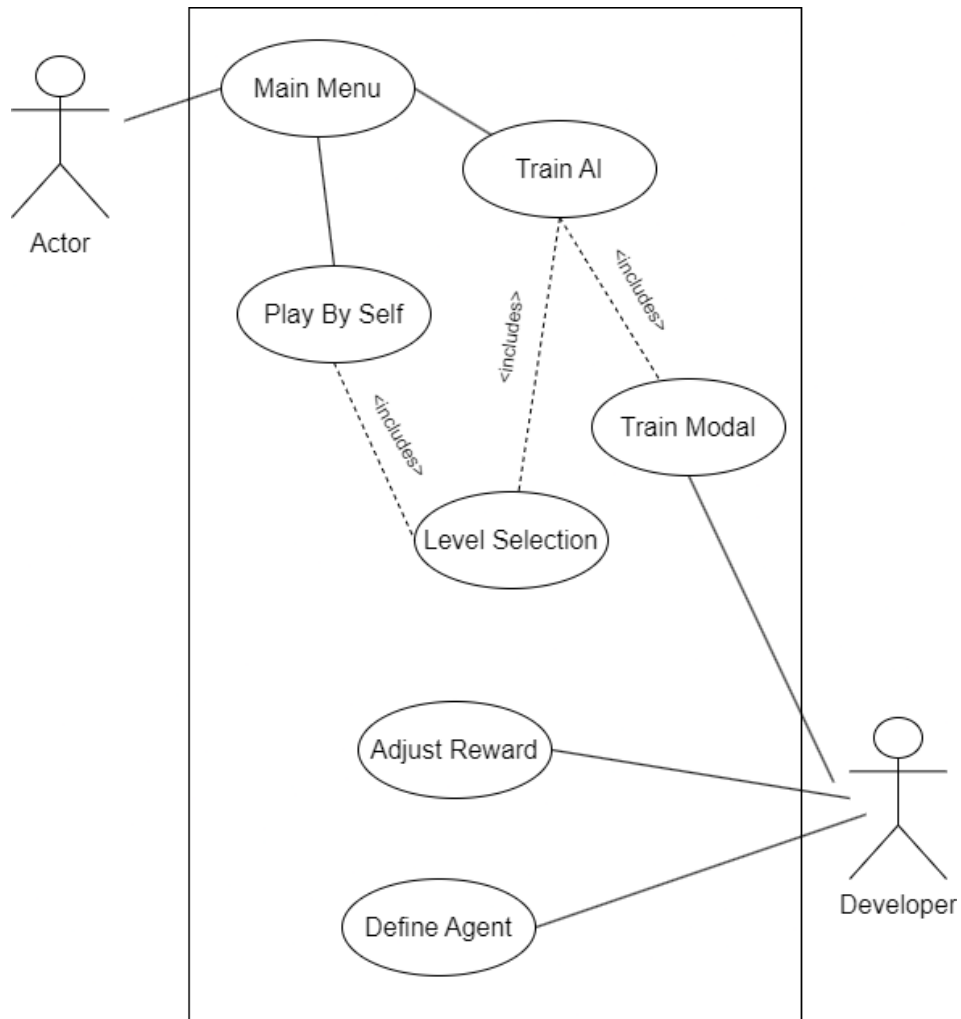
*Figure 3.5: Use Case Diagram*

### 3.2.4 Decision Tree

This decision tree represents the structured decision-making process used by the AI Agent in the "World's Hardest Game" project. It outlines how the AI evaluates different possible actions based on the current state of the game to maximize its chances of reaching the destination while avoiding enemies.
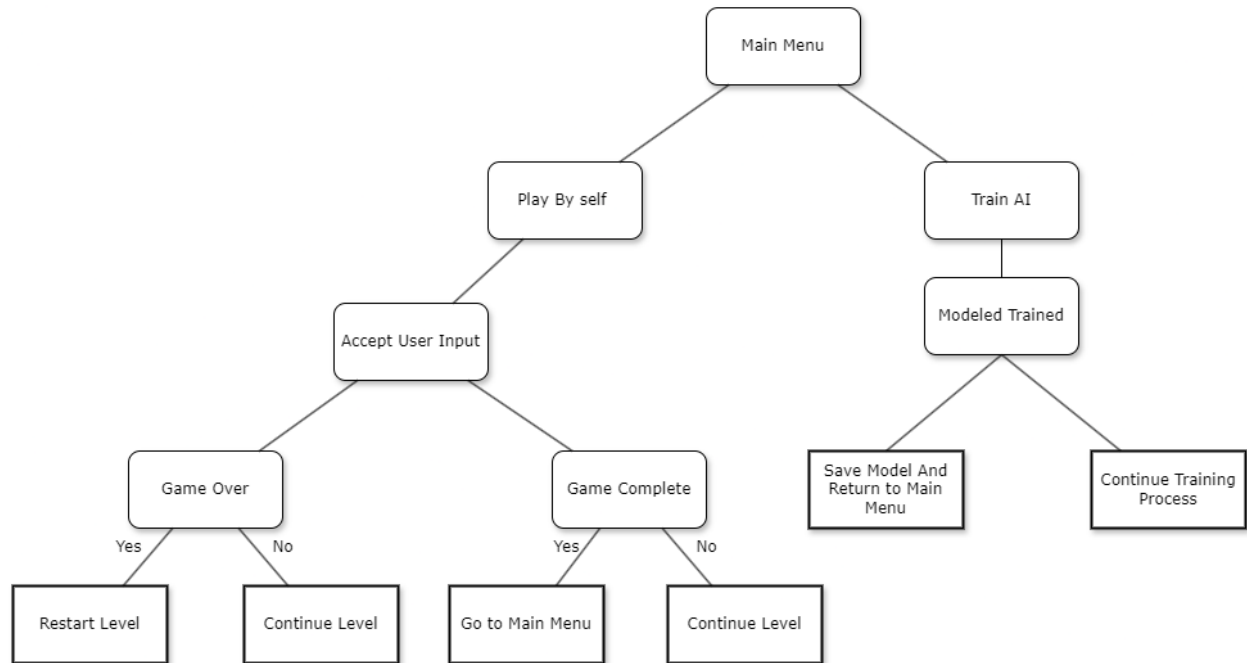
*Figure 3.6: Decision Tree*

### 3.2.5 Decision Table

This decision table outlines the decision-making criteria and actions for the AI Agent in the "World's Hardest Game" project. It specifies the possible states, conditions, and corresponding actions that the AI takes based on its environment and the reinforcement learning algorithm. The table is given below:

| CONDITIONS | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| UP Key Pressed | Y | N | N | N | Y | Y | N | N |
| Down Key Pressed | N | Y | N | N | N | N | Y | Y |
| Right Key Pressed | N | N | Y | N | Y | N | Y | N |
| Left Key Pressed | N | N | N | Y | N | Y | N | Y |

| ACTIONS | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| Move Player Up | ✓ | | | | ✓ | ✓ | | |
| Move Player Down | | ✓ | | | | | ✓ | ✓ |
| Move Player Right | | | ✓ | | ✓ | | ✓ | |
| Move Player Left | | | | ✓ | | ✓ | | ✓ |

*Figure 3.7: Decision Table*

### 3.2.6 Sequence Diagram

This sequence diagram illustrates the dynamic interaction between the main components of the AI system during a single game cycle in the "World's Hardest Game". It captures the sequential flow of actions performed by the Agent, GameAI, and the Reinforcement Learning Model (Linear_QNet and QTrainer) as the AI agent processes the current game state, decides on an action, and learns from the outcomes.

*Figure 3.8 Sequence Diagram*

**3.2.7 Class Diagram**

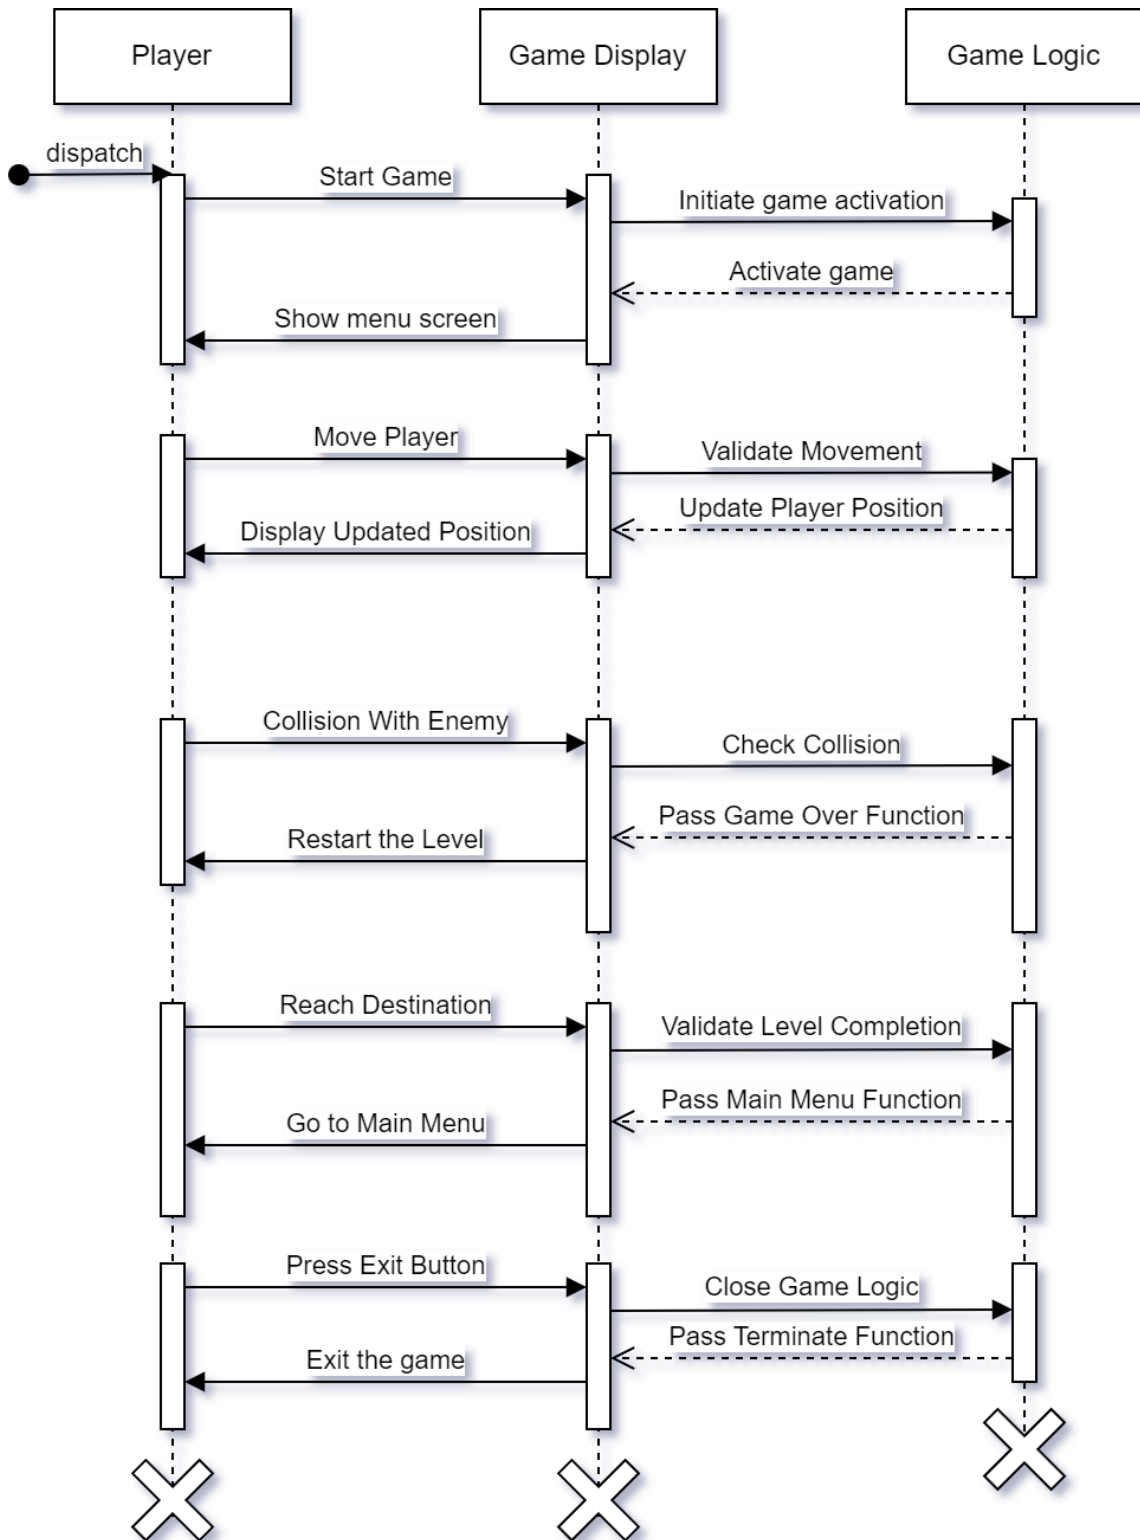This class diagram represents the static structure of the AI system implemented for the "World's Hardest Game". It includes the key classes and their relationships, providing a clear overview of the system's architecture and how various components interact with each other.
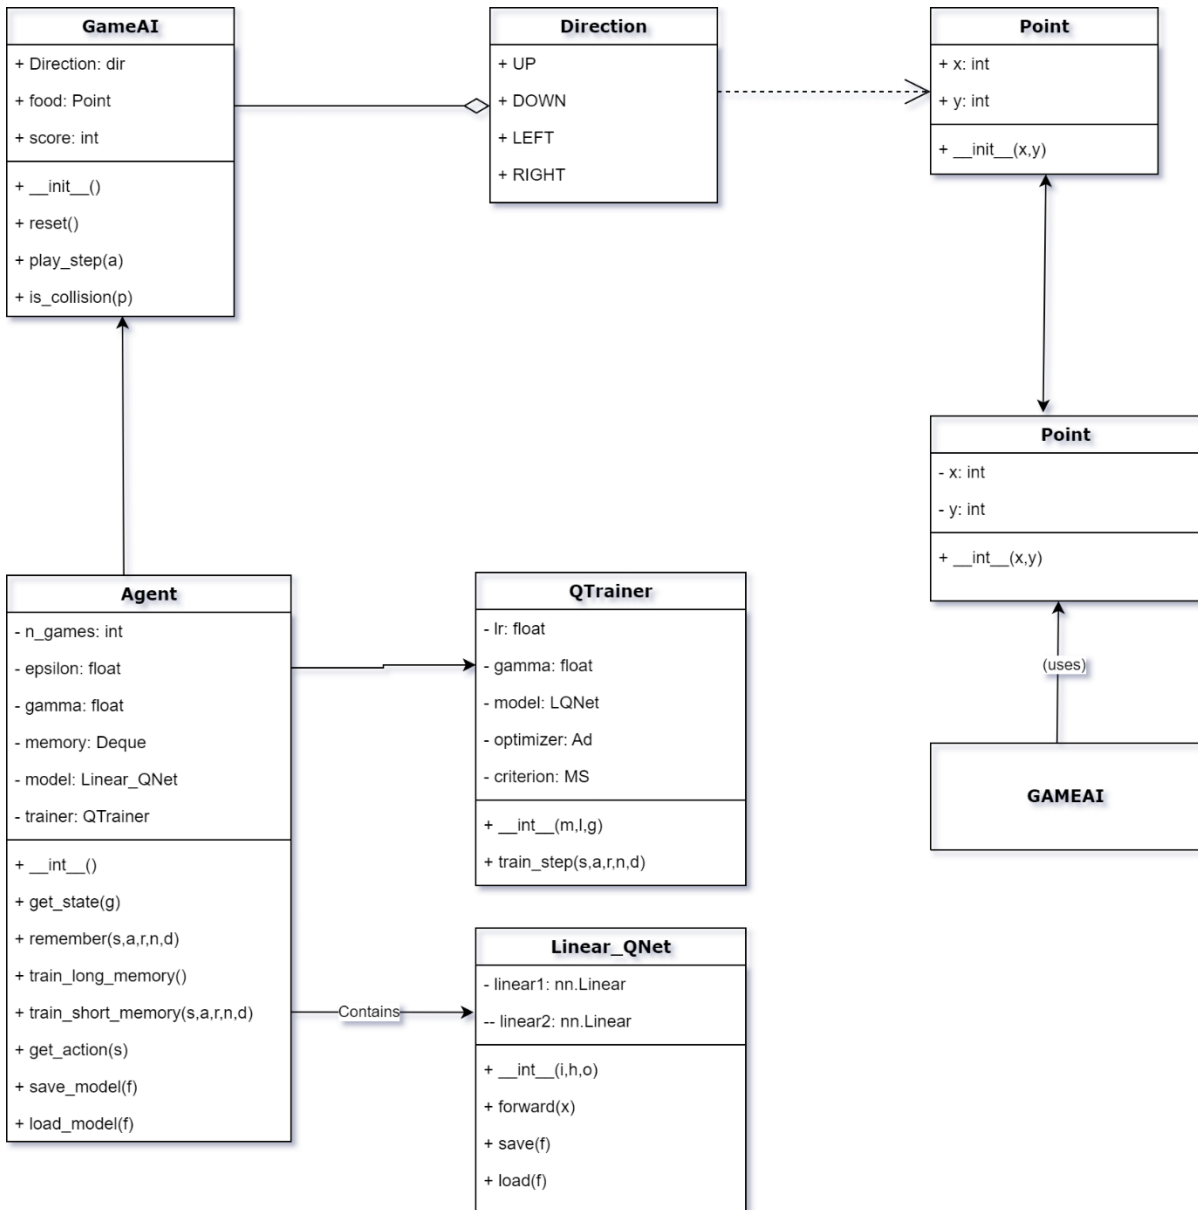


*Figure 3.9 Class Diagram*

## 3.3 Reinforcement Learning

The RL framework allows the agent to make informed decisions by learning from the rewards and penalties received during gameplay, ultimately improving its performance over time. The project leverages Reinforcement Learning (RL) to develop an intelligent agent capable of navigating through the "World's Hardest Game" by learning optimal strategies to avoid enemies and reach the destination.

### 3.3.1 Model Architecture

The RL agent is implemented using a Deep Q-Network (DQN) approach, which combines Q-learning with deep neural networks to approximate the Q-values for state-action pairs.

- Neural Network: The neural network consists of:
- Input Layer: Takes the state vector as input.
- Hidden Layers: Multiple fully connected layers with ReLU activations to capture complex features and interactions.
- Output Layer: Outputs Q-values for each possible action.
- Loss Function: Mean Squared Error (MSE) is used to minimize the difference between predicted Q-values and target Q-values.
- Optimizer: Adam optimizer is used to update the network weights based on the gradients computed from the loss.

*Figure 3.10: Q-Learning Modal Architecture*

### 3.3.3 Exploration-Exploitation Strategy

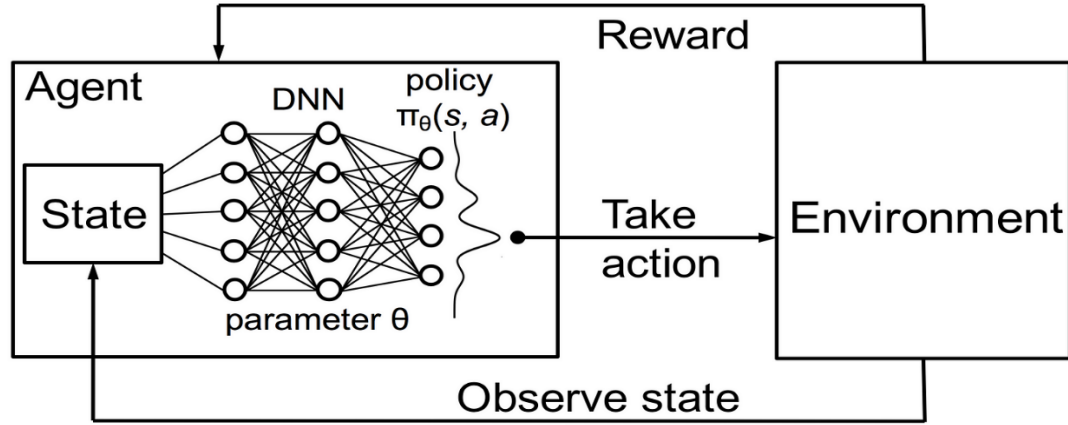To balance exploration of new strategies with exploitation of learned behaviors, an ε-greedy strategy is adopted. At each time step, the agent either selects a random action with probability ε (exploration) or chooses the action with the highest Q-value from the current state (exploitation). The exploration rate ε decays over time, promoting more exploitation as training progresses.

## 3.4 Training Process

The following section describes the modal training process.

1. **Experience Replay**: A memory buffer stores the agent's experiences (state, action, reward, next state, done) to break the correlation between consecutive experiences and improve training stability.

2. **Epsilon-Greedy Policy**: An exploration-exploitation strategy where the agent selects random actions with probability ε and the action with the highest Q-value with probability 1-ε. ε is decayed over time to shift from exploration to exploitation.

3. **Mini-Batch Training**: A mini-batch of experiences is sampled from the replay buffer, and the network is trained on this batch to update the Q-values.

4. **Target Network**: A separate target network is used to compute target Q-values, and its weights are periodically updated with the weights of the main network to stabilize training.

### 3.4.1 Training Loop

The processes of training loop are described below.

1. **State Extraction**: The agent extracts the current state from the game environment.
2. **Action Selection**: Based on the current state and ε-greedy strategy, the agent selects an action (move) to execute in the game.
3. **Environment Interaction**: The selected action is applied to the game environment, resulting in a reward and the next state.
4. **Memory Storage**: The tuple (state, action, reward, next state) is stored in the experience replay memory buffer.
5. **Batch Training**: Periodically, the agent samples a batch of experiences from the replay memory. For each batch, Q-values are computed using the DQN, and the loss is calculated to update the model parameters through backpropagation.
6. **Target Network Update**: To stabilize training, a target network periodically mirrors the current DQN weights.

The training phase involves optimizing the neural network's parameters using reinforcement learning techniques, enabling the AI agent to learn optimal strategies for navigating the 'World's Hardest Game' environment.
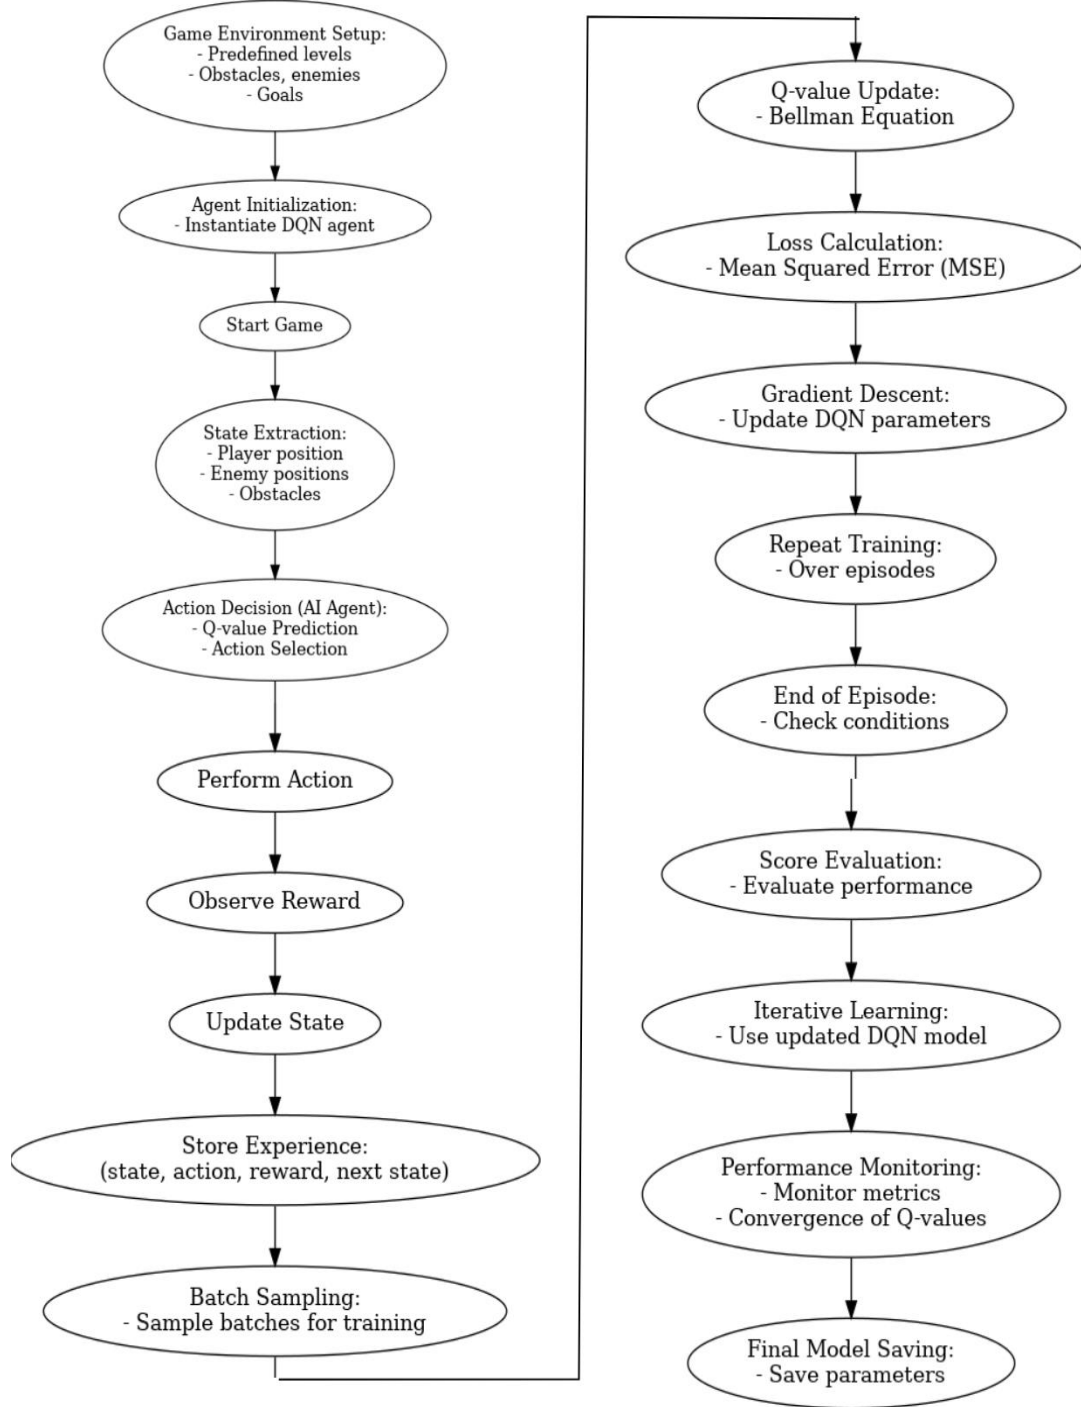
*Figure 3.11: Modal Training*

## 3.5 Evaluation

After the training of the model was completed, we evaluated the performance of the modal based on several metrics as discussed below:

### 3.5.1 Performance Metrics

The agent's performance is evaluated based on several metrics:

1.  **Score**: Total score accumulated across multiple episodes.
2.  **Completion Rate**: Percentage of successfully completed levels.
3.  **Training Stability**: Convergence of Q-values and policy over training epochs.
4.  **Learning Curve**: Plot of scores and mean scores over training episodes, illustrating learning progress.



*Figure 3.12: Model Training Graph*

### 3.5.2 Analysis

The agent's performance is compared against baseline approaches and human gameplay benchmarks where applicable. Comparative analysis provides insights into the effectiveness and efficiency of the RL-based approach in mastering the "World's Hardest Game."

## 3.6 Ethical Considerations

Ethical considerations include ensuring fair competition between human players and AI agents, maintaining transparency in AI decision-making processes, and addressing potential biases in model training data. The project adheres to ethical guidelines for AI research and gaming applications.

## 3.7 System Requirement Specifications

A System Requirements Specification (SRS) document provides a detailed description of a software system to be developed, laying out the functional and nonfunctional requirements. Here is a structure for an SRS document:

### 3.7.1. Software Specification

The entire RL framework, including the game environment, DQN model, Q-learning algorithm, and training loop, is implemented using Python and PyTorch. Libraries such as NumPy and matplotlib are utilized for numerical computations and visualization of training results.

1. Operating System: Windows 10 or Linux (Ubuntu 18.04 LTS recommended).
2. Python: Version 3.7 or higher.
3. Python Libraries:
   a. PyTorch 1.9.0 or higher
   b. NumPy
   c. Matplotlib
   d. Pygame
   e. TorchVision (for any image processing if required)
   f. Other standard libraries (os, random, collections, etc.)

4. Development Environment:

   a. Anaconda or Miniconda (recommended for managing Python environments)

   b. Jupyter Notebook (optional, for interactive testing and debugging)

   c. Integrated Development Environment (IDE) such as Visual Studio Code or PyCharm

5. Game Requirements:

   a. The "World's Hardest Game" or a similar game environment, playable through Pygame.

   b. Various levels with increasing difficulty, each containing obstacles, enemies, and goals.

6. AI Model Requirements:

   a. Implementation of a Q-learning based reinforcement learning model.

   b. Deep Q-Network (DQN) architecture or similar, implemented using PyTorch.

   c. Model capable of processing game state inputs and predicting Q-values for actions.

## 3.7.2 Hardware Specifications

The optimal hardware requirements are as follows:

1. Processor: Minimum Intel Core i5 or equivalent.

2. Memory: Minimum 8GB RAM.

3. Storage: Sufficient disk space to install necessary software and store game assets and trained AI models.

4. Graphics: A GPU with CUDA support (recommended for faster training times, optional but beneficial).

5. Display: Minimum resolution of 1280x720 pixels.

# Chapter 4: Discussion of the Achievements

The project marks a significant achievement with the successful integration of deep reinforcement learning techniques in the 'World's Hardest Game'. Key to this success is the effective use of the Q-learning algorithm, which enabled our AI agent to navigate complex game levels, evade enemies strategically, and achieve target objectives. Through rigorous training and optimization of neural network models, we achieved precise Q-values crucial for making informed decisions during gameplay. Our AI agent demonstrated adaptability across three diverse game levels, showcasing its ability to generalize learned behaviors and strategies across varied environments. Continuous refinement of our reinforcement learning model consistently led to high scores, occasionally outperforming human players in challenging scenarios. The integration of a user-friendly interface enhanced engagement for both human players and the AI agent, creating an immersive gaming experience. Looking forward, this project establishes a robust foundation for future advancements and scalability, indicating potential applications in more intricate gaming scenarios and real-world domains requiring sophisticated decision-making capabilities.

## 4.1 Overview

The "World's Hardest Game" project aimed to develop a robust reinforcement learning (RL) agent capable of autonomously playing a challenging game environment across three distinct levels. This section discusses the achievements, challenges encountered, and insights gained during the development and evaluation of the RL agent.

## 4.2 Achievement of Learning Objectives

Throughout the process of game development, model creation and development, several feats were achieved, and objectives were met. They have been described as follows:

### 4.2.1 Successful Creation of Game from scratch

The primary achievement of the project was the successful recreation of "World's Hardest Game." Three dynamic separate levels were created and implemented along with a minimalistic yet functional main menu and level selection screen. Some beautiful screenshots of various gameplay elements have been presented below:
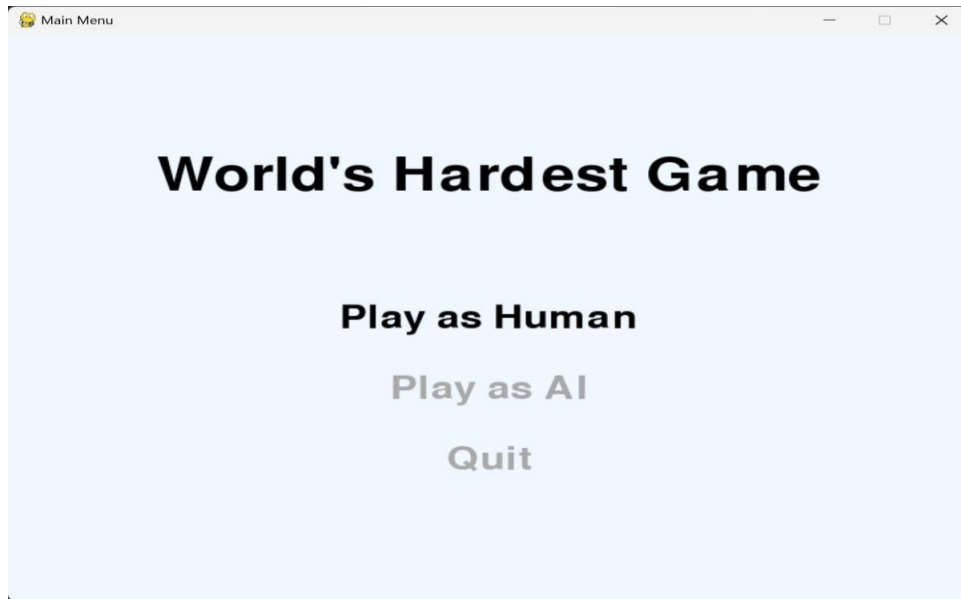
*Figure 4.1: Main Menu*



*Figure 4.2: Level 1*

*Figure 4.3: Level 2*
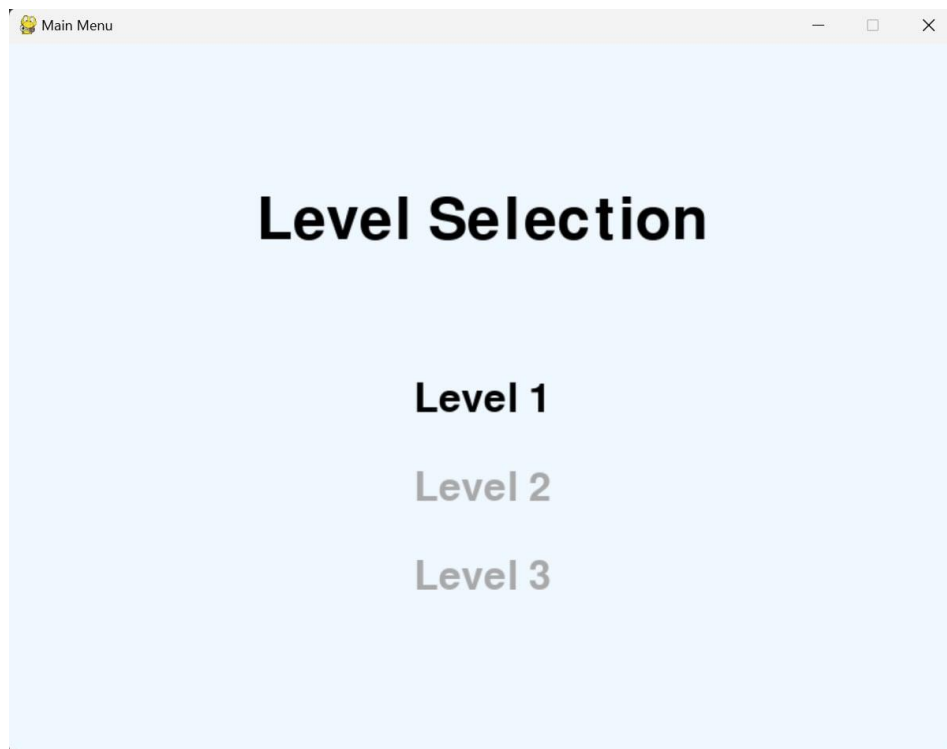
## 4.2.2 Successful Implementation of RL Framework

The prime achievement of the project also was the successful implementation of a reinforcement learning framework tailored for the "World's Hardest Game." The RL agent was equipped with Q-learning algorithm, enabling it to learn optimal policies through interaction with the game environment.



*Figure 4.4: Number of Games vs Score*

### 4.2.3 Autonomous Gameplay Across Multiple Levels

The RL agent demonstrated competence in navigating through two levels of increasing difficulty. It learned to avoid enemies, navigate around obstacles, and reach designated goals while optimizing its score. The agent's ability to generalize strategies across different level layouts showcased the effectiveness of the developed RL framework.

## 4.3 Performance Evaluation

The performance of RL agent was done as follows:

### 4.3.1 Score Accumulation and Training Stability

During training, the RL agent exhibited a steady increase in cumulative scores across episodes. The learning curve illustrated progressive improvement in the agent's ability to maximize rewards and avoid penalties. Stable convergence of Q-values and policy was observed, indicating effective learning and decision-making capabilities.
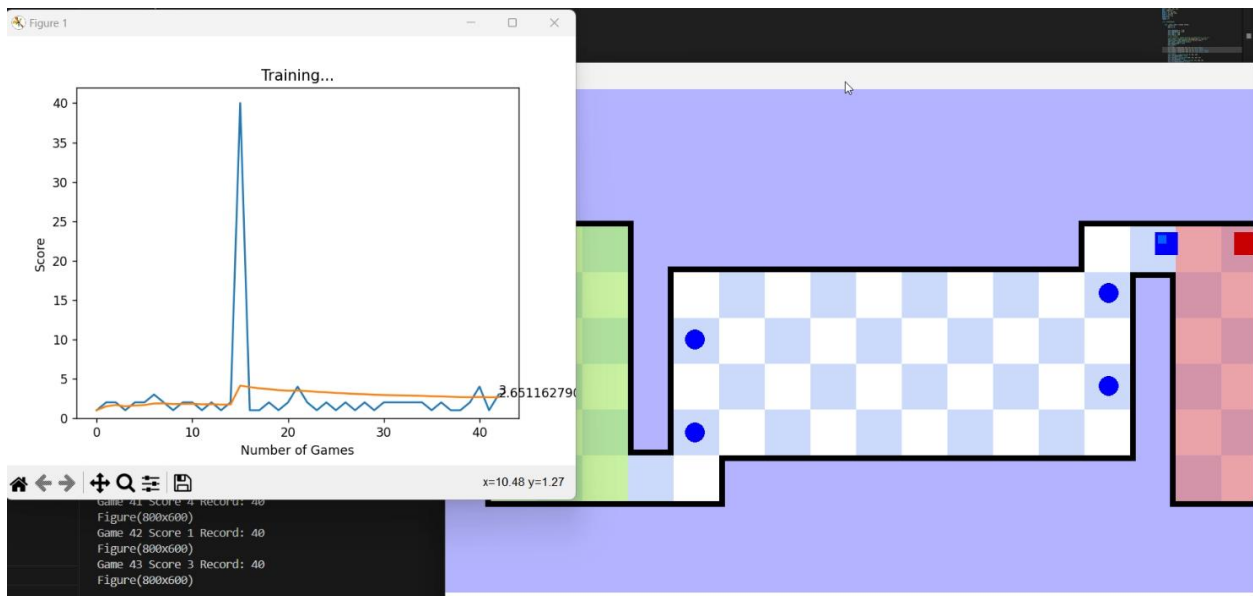


*Figure 4.5: Score Accumulation and Training Stability*

### 4.3.2 Completion Rate and Efficiency

The RL agent achieved a high completion rate in successfully navigating through the game levels. It demonstrated efficiency in learning optimal trajectories, leveraging the Q-learning algorithm to prioritize actions that lead to goal attainment while avoiding collisions with enemies.

## 4.4 Comparative Analysis

The comparisons of Reinforcement Learning agent were done as follows:

### 4.4.1 Comparison with Baseline Approaches

The performance of the RL agent was compared against baseline approaches, including random policy agents or heuristics-based strategies. The results highlighted significant improvements in score accumulation and completion rate, underscoring the superiority of RL-based methods in mastering complex gameplay environments.
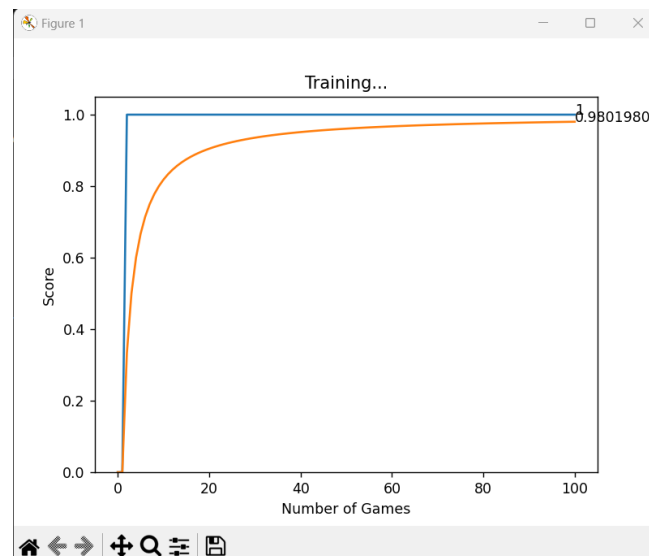


*Figure 4.6 Learning rate using random policy agents*

Using random policy agents, the agent could barely reach the initial checkpoint even after 100 games. The RL approach was much more efficient , reaching the fourth checkpoint in just over 35 games as shown in the figure below:

*Figure 4.7: Learning rate using RL approach*

### 4.4.2 Human Benchmarking

Where applicable, the RL agent's performance was benchmarked against human gameplay data or expert strategies. While human players may initially outperform the agent due to intuitive decision-making and adaptability, once the model was trained, the RL agent demonstrated consistent performance and stability in solving the level.

## 4.5 Challenges and Limitations

The completion of this project was not an easy feat to achieve. We were bounded by several challenges and limitations.

### 4.5.1 Training Convergence and Hyperparameter Tuning

One of the primary challenges encountered was achieving stable training convergence. Fine-tuning hyperparameters such as learning rate, exploration rate decay, and batch size was crucial in balancing exploration-exploitation trade-offs and optimizing learning efficiency.

### 4.5.2 Generalization Across Levels

Generalizing learned policies across diverse level configurations posed another challenge. While the RL agent adapted well to varying layouts of obstacles and enemies, occasional performance discrepancies were observed in overly complex scenarios requiring more nuanced decision-making.

## 4.6 Insights and Future Directions

This section covers insights and future enhancements to further improve learning stability and efficiency.

### 4.6.1 Insights from RL Agent Behavior

Insights gained from the RL agent's behavior included the emergence of strategic maneuvers such as pathfinding around enemies and prioritizing shorter routes to goals. Observations of exploration-exploitation dynamics provided valuable insights into reinforcement learning dynamics applicable beyond gaming contexts.

### 4.6.2 Future Directions for Enhancement

Future enhancements could focus on integrating advanced RL techniques such as Double Q-learning or Dueling DQN architectures to further improve learning stability and efficiency. Enhancing the agent's robustness to novel environments and dynamic obstacles would expand its applicability to real-world scenarios.

## 4.7 Ethical Considerations

Ethical considerations encompassed ensuring fair competition between RL agents and human players, transparent communication of AI decision-making processes, and addressing biases in training data. Upholding ethical standards in AI development and deployment remains essential for fostering trust and accountability in gaming and AI research.

# Chapter 5: Conclusion and Recommendations

The development and implementation of a reinforcement learning (RL) agent for playing the "World's Hardest Game" have proven to be a significant success. The project successfully demonstrated the capability of an RL agent to autonomously navigate through complex game environments, avoiding enemies and reaching goals across three distinct levels.

## 5.1 Key Findings

The key findings of our projects include the followings:

1. **Effective Learning Framework**: The integration of Q-learning algorithm facilitated robust learning of optimal policies, enabling the agent to maximize rewards while minimizing penalties.
2. **Performance and Efficiency**: The RL agent exhibited commendable performance in terms of score accumulation, completion rates, and decision-making efficiency. It showcased the potential of RL in mastering intricate gameplay dynamics.
3. **Comparative Analysis**: Comparative assessments against baseline approaches and human benchmarks underscored the superiority of RL-based strategies in achieving competitive gameplay outcomes.

## 5.2 Contributions to Knowledge

The project contributed valuable insights into the application of RL techniques in gaming contexts. It highlighted the efficacy of reinforcement learning algorithms in enhancing autonomous decision-making and strategic planning within dynamic and challenging environments.

## 5.3 Recommendations for Future Work

The project can be improved and re-iterated to further improve and enhance for higher level of implications. Such recommendations are mentioned as follows:

### 5.3.1 Advanced RL Techniques

Future research could explore advanced RL methodologies such as Double Q-learning, Prioritized Experience Replay, or Actor-Critic architectures. These techniques could further enhance the agent's learning stability, convergence speed, and adaptability to novel game scenarios.

### 5.3.2 Real-World Applications

Expanding the RL agent's applicability beyond gaming to real-world applications could be a promising avenue. Examples include autonomous navigation in dynamic environments, robotic control systems, and adaptive decision-making in industrial processes.

### 5.3.3 Generalization and Transfer Learning

Improving the agent's ability to generalize learned policies across diverse game levels and environments remains a critical area of development. Techniques like transfer learning and domain adaptation could facilitate faster adaptation to new challenges and scenarios.

### 5.3.4 Ethical Considerations

Continued attention to ethical considerations in AI development is essential. Ensuring fairness, transparency in decision-making, and mitigating biases in training data are imperative for responsible deployment of AI technologies in gaming and beyond.

## 5.4 Conclusion Statement

In conclusion, the project successfully designed and implemented a reinforcement learning framework for the "World's Hardest Game," achieving competitive performance and advancing knowledge in autonomous gameplay strategies. The insights gained pave the way for future advancements in AI research and applications across various domains.

# References

*AlphaGo*. (n.d.). Retrieved June 16, 2024, from Google DeepMind: https://deepmind.google/technologies/alphago/?fbclid=IwZXh0bgNhZW0CMTAAAR1n 5AS762paE4xqYzyntHOsF_q7U_po7ofc8nCE- _c57zd4nJVW7X8mBog_aem_ZmFrZWR1bW15MTZieXRlcw

bajaj, p. (2023, April 18). *Reinforcement Learning*. Retrieved June 16, 2024, from geeksforgeeks: https://www.geeksforgeeks.org/what-is-reinforcement-learning/

Chanda, K. K. (2024, March 21). *Q-Learning in Python*. Retrieved June 16, 2024, from GeeksforGeeks: https://www.geeksforgeeks.org/q-learning-in-python/

CodeBullet. (2018, April 29). Retrieved from YouTube: https://www.youtube.com/watch?v=sB_IGstiWlc&ab_channel=CodeBullet

CoderBoi. (2022, October 22). Retrieved from YouTube: https://www.youtube.com/watch?v=L2usRyiB9Vs

Jain, S. (2023, April 18). *Reinforcement learning*. Retrieved June 15, 2024, from GeeksforGeeks: https://www.geeksforgeeks.org/what-is-reinforcement-learning/

Julian Togelius, N. S. (2013, September 01). *The Mario AI Championship*. Retrieved June 16, 2024, from Wiley: https://onlinelibrary.wiley.com/doi/10.1609/aimag.v34i3.2492?fbclid=IwZXh0bgNhZW0 CMTAAAR3vB44qcs5sQ5tKaUCkfGOW_kJC_lhK4fAZRrhHqew7XaVgA2kUIjFyDak _aem_ZmFrZWR1bW15MTZieXRlcw

Oscar, A. (2023, October 19). *Generative AI in Game Design*. Retrieved June 16, 2024, from medium: https://medium.com/@aishwarya_62914/generative-ai-in-game-design- procedural-content-generation-and-ai-driven-npcs-d93dc15d39d9

Precup, D. (n.d.). *Doina precup*. Retrieved June 16, 2024, from Mila: https://mila.quebec/en/doina-precup

team, T. A. (2019, January 24). *AlphaStar*. Retrieved June 16, 2024, from Google DeepMind: https://deepmind.google/discover/blog/alphastar-mastering-the-real-time-strategy-game-starcraft-ii/?fbclid=IwZXh0bgNhZW0CMTAAAR1FQB9hBwPsNKUTjHzVOz_WWuaRjTZrSpkJy0m0dP9-PLhZTPlEh75M9Iw_aem_ZmFrZWR1bW15MTZieXRlcw

*Worlds Hardest Game*. (n.d.). Retrieved June 16, 2024, from carzygames.com: https://www.crazygames.com/game/worlds-hardest-game